

Autoencoders in Machine Learning

Last Updated : 09 Oct, 2025

Autoencoders are a special type of neural networks that learn to compress data into a compact form and then reconstruct it to closely match the original input. They consist of an:

- Encoder that captures important features by reducing dimensionality.
- Decoder that rebuilds the data from this compressed representation.

The model trains by minimizing reconstruction error using loss functions like Mean Squared Error or Binary Cross-Entropy. These are applied in tasks such as noise removal, error detection and feature extraction where capturing efficient data representations is important.

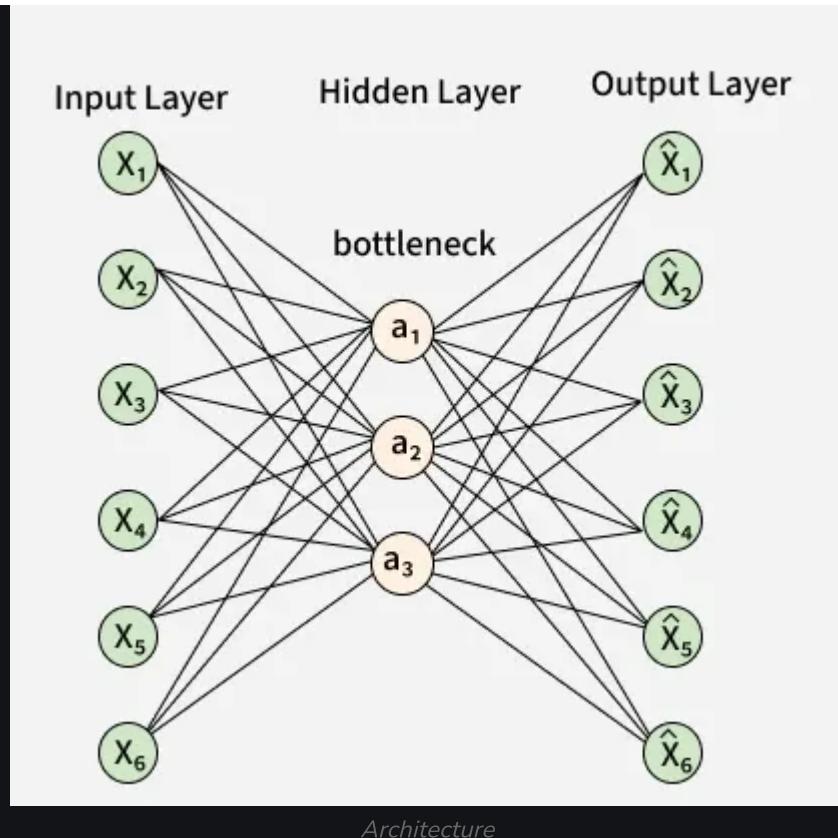
What is an Autoencoder?

- Learns to compress and reconstruct data
- Two parts: Encoder (compress), Decoder (reconstruct)
- Used for noise removal, anomaly detection & feature extraction

The diagram illustrates the architecture of an autoencoder. It shows three main layers: an 'Encoder' at the top, a 'Latent Space' in the middle, and a 'Decoder' at the bottom. The 'Input Data' layer on the left has arrows pointing upwards to the 'Encoder' and downwards to the 'Latent Space'. The 'Encoder' layer has arrows pointing downwards to the 'Latent Space' and upwards to the 'Decoder'. The 'Latent Space' layer has arrows pointing upwards to the 'Decoder' and downwards to the 'Reconstructed Data' layer on the right. The 'Decoder' layer has arrows pointing downwards to the 'Reconstructed Data' layer and upwards to the 'Latent Space'.

Architecture of Autoencoder

An autoencoder's architecture consists of three main components that work together to compress and then reconstruct data which are as follows:



Architecture

1. Encoder

It compresses the input data into a smaller, more manageable form by reducing its dimensionality while preserving important information. It has three layers which are:

- **Input Layer:** This is where the original data enters the network. It can be images, text features or any other structured data.
- **Hidden Layers:** These layers perform a series of transformations on the input data. Each hidden layer applies weights and [activation functions](#) to capture important patterns, progressively reducing the data's size and complexity.
- **Output(Latent Space):** The encoder outputs a compressed vector known as the latent representation or encoding. This vector captures the important features of the input data in a condensed form helps in filtering out noise and redundancies.

2. Bottleneck (Latent Space)

It is the smallest layer of the network which represents the most compressed version of the input data. It serves as the information

bottleneck which force the network to prioritize the most significant features. This compact representation helps the model learn the underlying structure and key patterns of the input helps in enabling better generalization and efficient data encoding.

3. Decoder

It is responsible for taking the compressed representation from the latent space and reconstructing it back into the original data form.

- **Hidden Layers:** These layers progressively expand the latent vector back into a higher-dimensional space. Through successive transformations decoder attempts to restore the original data shape and details
- **Output Layer:** The final layer produces the reconstructed output which aims to closely resemble the original input. The quality of reconstruction depends on how well the encoder-decoder pair can minimize the difference between the input and output during training.

Loss Function in Autoencoder Training

During training an autoencoder's goal is to minimize the reconstruction loss which measures how different the reconstructed output is from the original input. The choice of loss function depends on the type of data being processed:

- **Mean Squared Error (MSE):** This is commonly used for continuous data. It measures the average squared differences between the input and the reconstructed data.
- **Binary Cross-Entropy:** Used for binary data (0 or 1 values). It calculates the difference in probability between the original and reconstructed output.

During training the network updates its weights using [backpropagation](#) to minimize this reconstruction loss. By doing this it learns to extract and retain the most important features of the input data which are encoded in the latent space.

Efficient Representations in Autoencoders

Constraining an autoencoder helps it learn meaningful and compact features from the input data which leads to more efficient representations. After training only the encoder part is used to encode similar data for future tasks. Various techniques are used to achieve this are as follows:

- **Keep Small Hidden Layers:** Limiting the size of each hidden layer forces the network to focus on the most important features. Smaller layers reduce redundancy and allows efficient encoding.
- **Regularization:** Techniques like [L1 or L2 regularization](#) add penalty terms to the loss function. This prevents overfitting by removing excessively large weights which helps in ensuring the model to learns general and useful representations.
- **Denoising:** In denoising autoencoders random noise is added to the input during training. It learns to remove this noise during reconstruction which helps it focus on core, noise-free features and helps in improving robustness.
- **Tuning the Activation Functions:** Adjusting activation functions can promote sparsity by activating only a few neurons at a time. This sparsity reduces model complexity and forces the network to capture only the most relevant features.

Types of Autoencoders

Lets see different types of Autoencoders which are designed for specific tasks with unique features:

1. Denoising Autoencoder

[Denoising Autoencoder](#) is trained to handle corrupted or noisy inputs, it learns to remove noise and helps in reconstructing clean data. It prevent the network from simply memorizing the input and encourages learning the core features.

2. Sparse Autoencoder

[Sparse Autoencoder](#) contains more hidden units than input features but only allows a few neurons to be active simultaneously. This sparsity is controlled by zeroing some hidden units, adjusting activation functions or adding a sparsity penalty to the loss function.

3. Variational Autoencoder

[Variational autoencoder \(VAE\)](#) makes assumptions about the probability distribution of the data and tries to learn a better approximation of it. It uses [stochastic gradient descent](#) to optimize and learn the distribution of latent variables. They are used for generating new data such as creating realistic images or text.

It assumes that the data is generated by a Directed Graphical Model and tries to learn an approximation to $q_\phi(z|x)$ to the conditional property $q_\theta(z|x)$ where ϕ and θ are the parameters of the encoder and the decoder respectively.

4. Convolutional Autoencoder

Python for Machine Learning Machine Learning with R Machine Learning Algorithms EDA 

(CNNs) which are designed for processing images. The encoder extracts features using convolutional layers and the decoder reconstructs the image through **deconvolution** also called as upsampling.

Implementation of Autoencoders

We will create a simple autoencoder with two Dense layers: an encoder that compresses images into a 64-dimensional latent vector and a decoder that reconstructs the original image from this compressed form.

Step 1: Import necessary libraries

We will be using [Matplotlib](#), [NumPy](#), [TensorFlow](#) and the MNIST dataset loader for this.

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, losses
from tensorflow.keras.models import Model
from keras.datasets import mnist
```

Step 2: Load the MNIST dataset

We will be loading the MNIST dataset which is inbuilt dataset and normalize pixel values to [0,1] also reshape the data to fit the model.

```
(x_train, _), (x_test, _) = mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

x_train = np.reshape(x_train, (len(x_train), 28, 28, 1))
x_test = np.reshape(x_test, (len(x_test), 28, 28, 1))
```

Output:

*Shape of the training data: (60000, 28, 28)
Shape of the testing data: (10000, 28, 28)*

Step 3: Define a basic Autoencoder

Creating a simple autoencoder class with an encoder and decoder using [Keras Sequential](#) model.

- **layers.Input(shape=(28, 28, 1))**: Input layer expecting grayscale images of size 28x28.
- **layers.Dense(latent_dimensions, activation='relu')**: Dense layer that compresses the input to the latent space using [ReLU](#) activation.
- **layers.Dense(28 * 28, activation='sigmoid')**: Dense layer that expands the latent vector back to the original image size with [sigmoid](#) activation.

```
class SimpleAutoencoder(Model):
    def __init__(self, latent_dimensions):
        super(SimpleAutoencoder, self).__init__()
        self.encoder = tf.keras.Sequential([
            layers.Input(shape=(28, 28, 1)),
```

```

        layers.Flatten(),
        layers.Dense(latent_dimensions, activation='relu'),
    ])

    self.decoder = tf.keras.Sequential([
        layers.Dense(28 * 28, activation='sigmoid'),
        layers.Reshape((28, 28, 1))
    ])

def call(self, input_data):
    encoded = self.encoder(input_data)
    decoded = self.decoder(encoded)
    return decoded

```

Step 4: Compiling and Fitting Autoencoder

Here we compile the model using [Adam optimizer](#) and [Mean Squared Error](#) loss also we train for 10 epochs with batch size 256.

- **latent_dimensions = 64:** Sets the size of the compressed latent space to 64.

```

latent_dimensions = 64
autoencoder = SimpleAutoencoder(latent_dimensions)
autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())

autoencoder.fit(x_train, x_train,
                 epochs=10,
                 batch_size=256,
                 shuffle=True,
                 validation_data=(x_test, x_test))

```

Output:

```

Epoch 1/10
235/235 3s 10ms/step - loss: 0.0973 - val_loss: 0.0321
Epoch 2/10
235/235 2s 9ms/step - loss: 0.0289 - val_loss: 0.0209
Epoch 3/10
235/235 2s 8ms/step - loss: 0.0197 - val_loss: 0.0151
Epoch 4/10
235/235 3s 10ms/step - loss: 0.0144 - val_loss: 0.0115
Epoch 5/10
235/235 3s 10ms/step - loss: 0.0112 - val_loss: 0.0093
Epoch 6/10
235/235 2s 9ms/step - loss: 0.0091 - val_loss: 0.0077
Epoch 7/10
235/235 2s 8ms/step - loss: 0.0076 - val_loss: 0.0067
Epoch 8/10
235/235 2s 8ms/step - loss: 0.0066 - val_loss: 0.0059
Epoch 9/10
235/235 3s 9ms/step - loss: 0.0060 - val_loss: 0.0055
Epoch 10/10
235/235 3s 10ms/step - loss: 0.0056 - val_loss: 0.0051

```

Training

Step 5: Visualize original and reconstructed data

Now compare original images and their reconstructions from the autoencoder.

- **encoded_imgs = autoencoder.encoder(x_test).numpy()**: Passes test images through the encoder to get their compressed latent representations as NumPy arrays.
- **decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()**: Reconstructs images by passing the latent representations through the decoder and converts them to NumPy arrays.

```

encoded_imgs = autoencoder.encoder(x_test).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

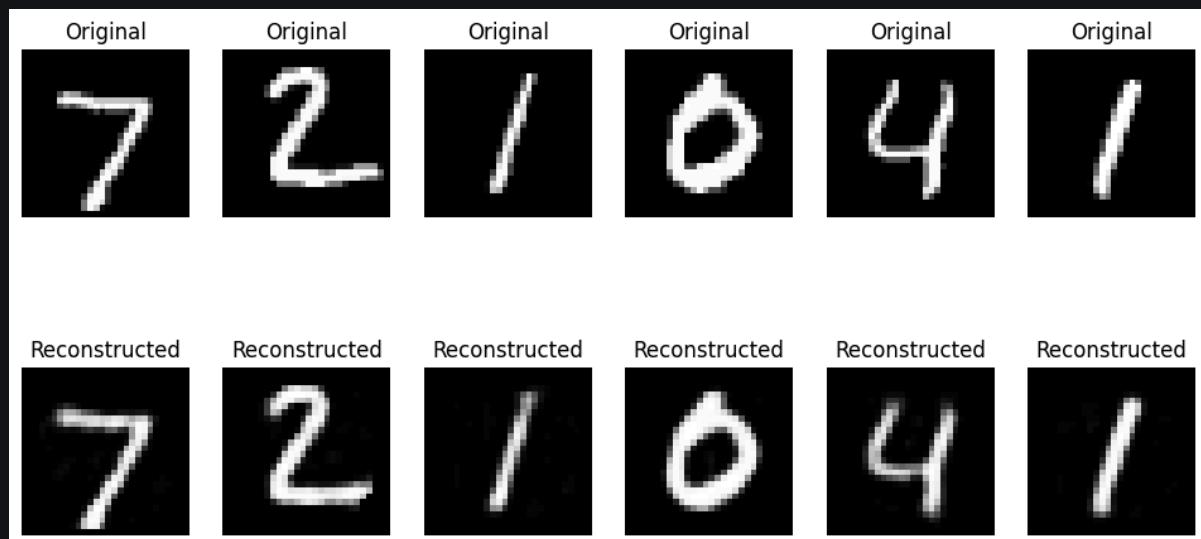
n = 6
plt.figure(figsize=(12, 6))
for i in range(n):
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    plt.title("Original")
    plt.axis('off')

    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
    plt.title("Reconstructed")
    plt.axis('off')

plt.show()

```

Output:



The visualization compares original MNIST images (top row) with their reconstructed versions (bottom row) showing that the autoencoder effectively captures key features despite some minor blurriness.

Limitations

Autoencoders are useful but also have some limitations:

- **Memorizing Instead of Learning Patterns:** It can sometimes memorize the training data rather than learning meaningful patterns which reduces their ability to generalize to new data.
- **Reconstructed Data Might Not Be Perfect:** Output may be blurry or distorted with noisy inputs or if the model architecture lacks sufficient complexity to capture all details.
- **Requires a Large Dataset and Good Parameter Tuning:** It requires large amounts of data and careful parameter tuning (latent dimension size, learning rate, etc) to perform well. Insufficient data or poor tuning can result in weak feature representations.

Comment



AlindG...
Follow

12

Article Tags :

Machine Learning

AI-ML-DS

AI-ML-DS With Python